



Don't Wait For Sync To Achieve Strong Consistency!

Seo Jin Park and John Ousterhout




OVERVIEW

Deterministic updates **durable** and **consistent** with **asynchronous** backup

- Server returns to clients before making updates durable.
- Use “retries” of RPCs to recover from server crash.
- Two ways to recover from crash.
 - Client-driven retries.
 - 3rd party (witness) driven retries.
- Better Performance
 - Write latency: 15 μ s \rightarrow 7.5 μ s
 - Better throughput: 3x server throughput

LET'S BACKUP ASYNC



asynchrony = weak consistency??

State of art: Consistency **OR** Performance

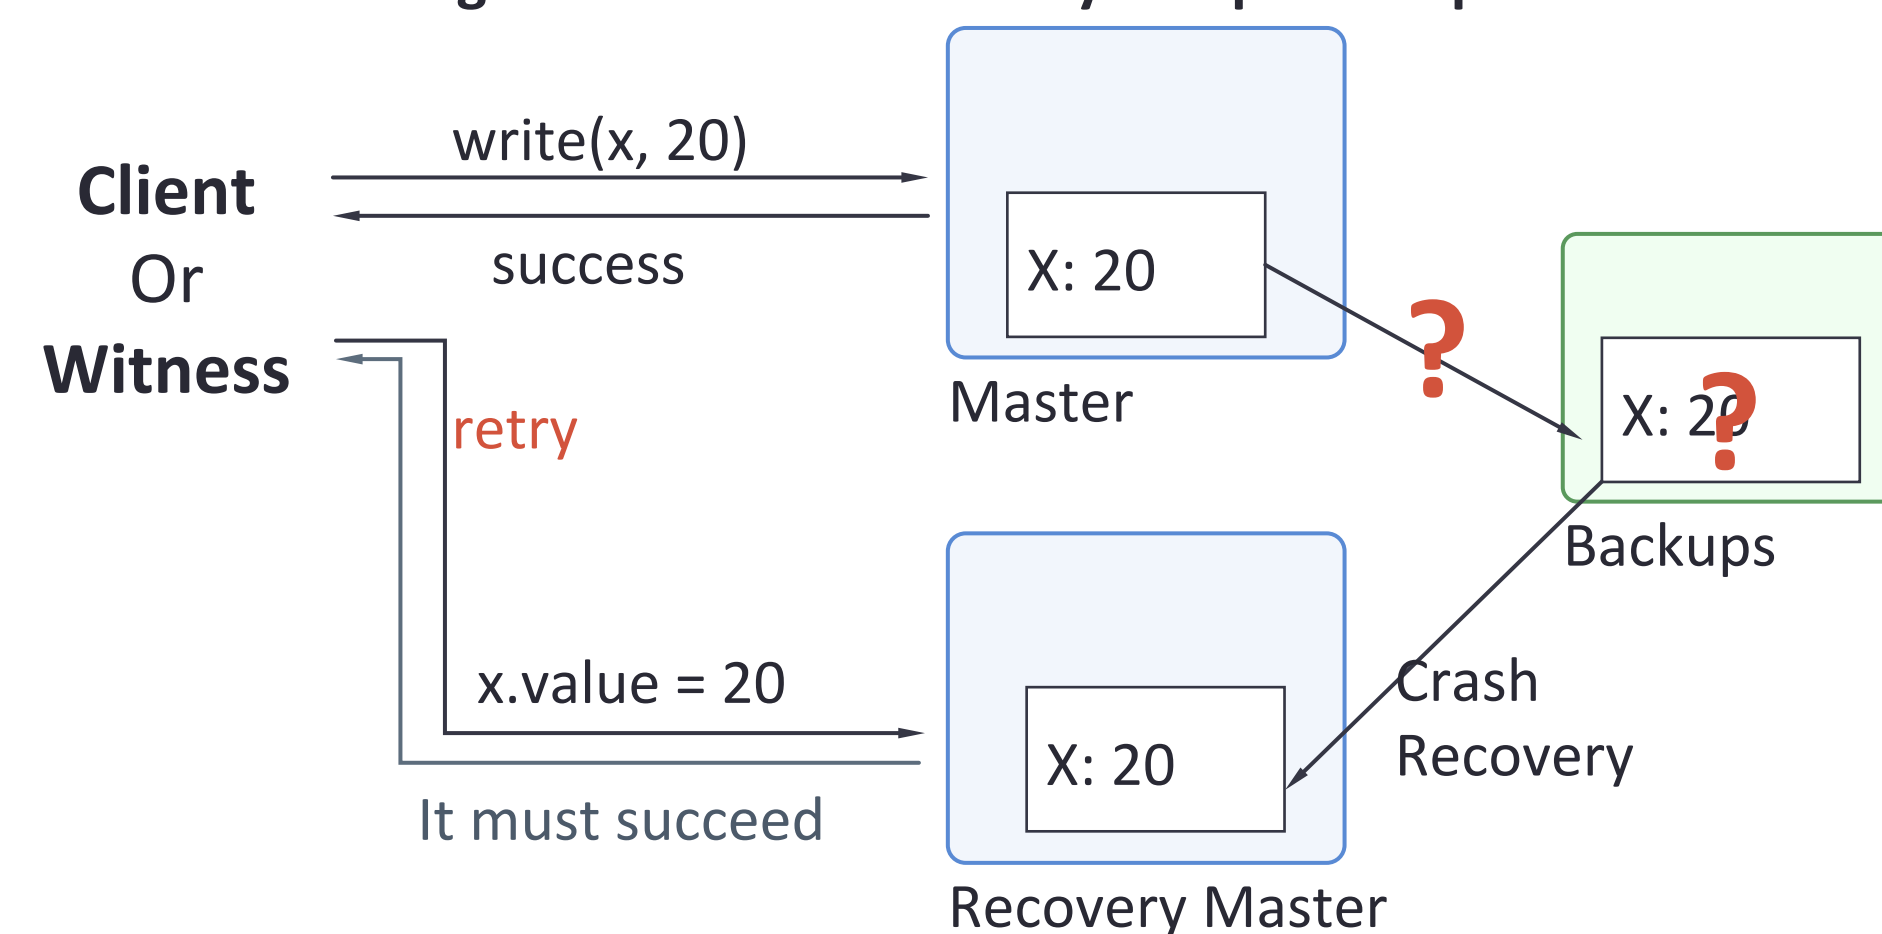
- Async replication + eventual consistency (eg. TAO, Redis)
- Sync replication + strong consistency (eg. RAMCloud)

We can get: Performance & Strong Consistency & simplicity of Primary-Backup

CHALLENGES IN RECOVER BY RETRY

Issue 1. Retry may re-execute

- If a server crash, an update may or may not be recovered
- RIFL (Reusable Infrastructure for Linearizability) [SOSP15] will let server ignore retries for already completed updates



Issue 2. Out-of-order retries

- Allow 1 not-replicated update per key (overwrites wait for sync of previous value)
 - Pro: Any deterministic operations can be recovered by retries.
 - Con: continuously overwritten object can be slow.

CLIENT-DRIVEN RECOVERY

New Consistency Model: consistency w/o durability

1. All reads are consistent
 - Reads are blocked until data become durable
 2. Client written data should never experience anomaly as long as it is alive.
 - When a server crashes, client retries previously returned writes.
- Write (last two) is lost only if both client and server crash
 - Client may wait for durability before externalization
 - Conditional write is still consistent
 - Non-durable atomic (ACI) transaction

Example

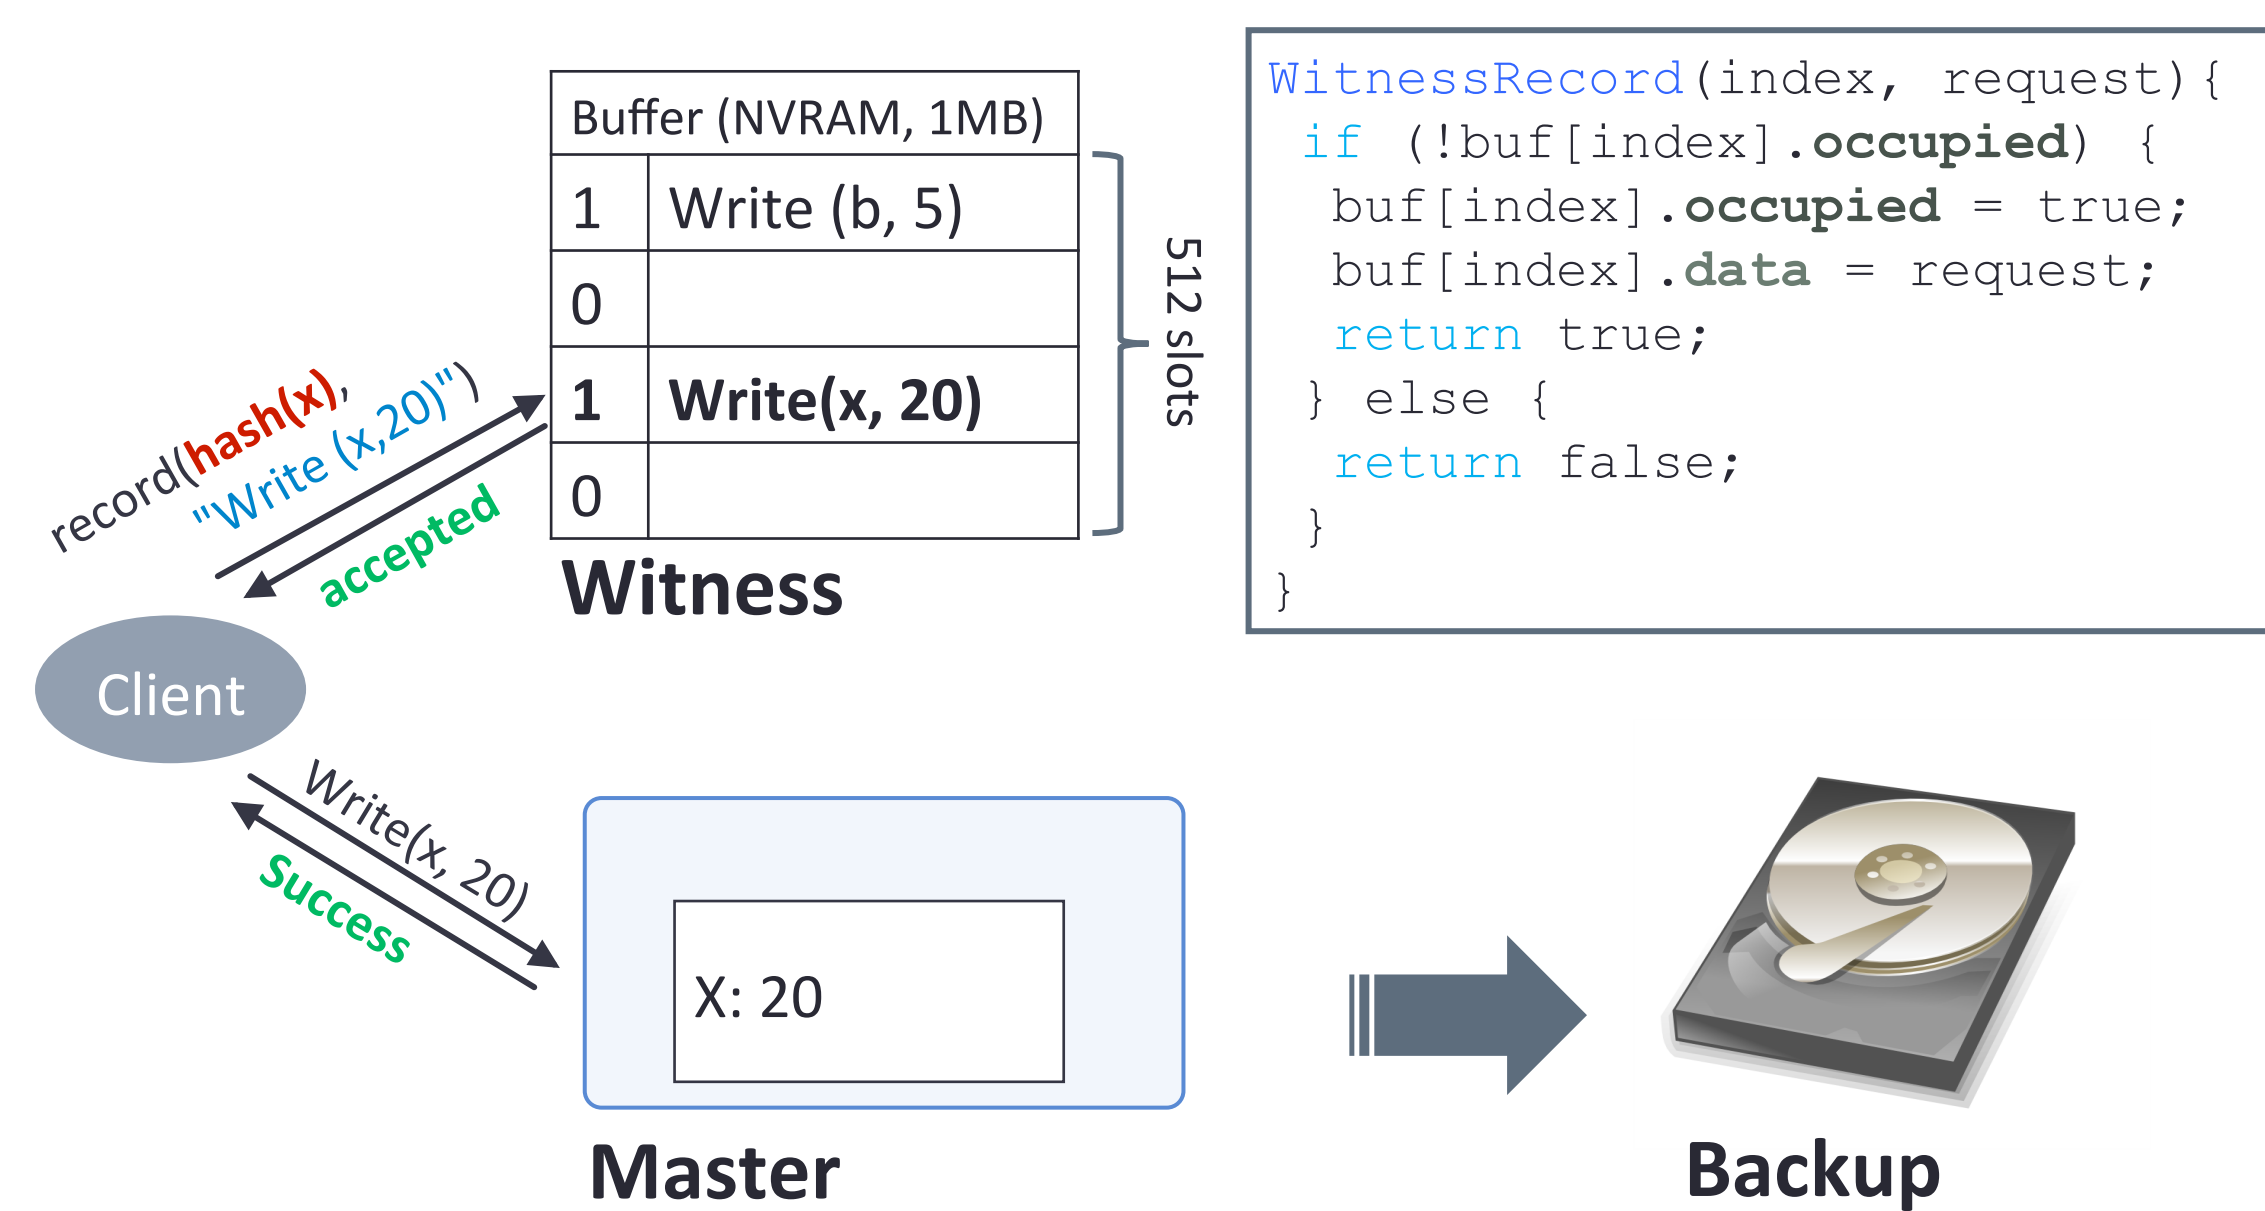
```
ramcloud.write(1, "Bob", "2");
ramcloud.write(1, "Bill", "2");
ramcloud.sync();
printf("Updated Bob and Bill");
```

Don't guarantee invariants. App still need to clean up for crashed clients.

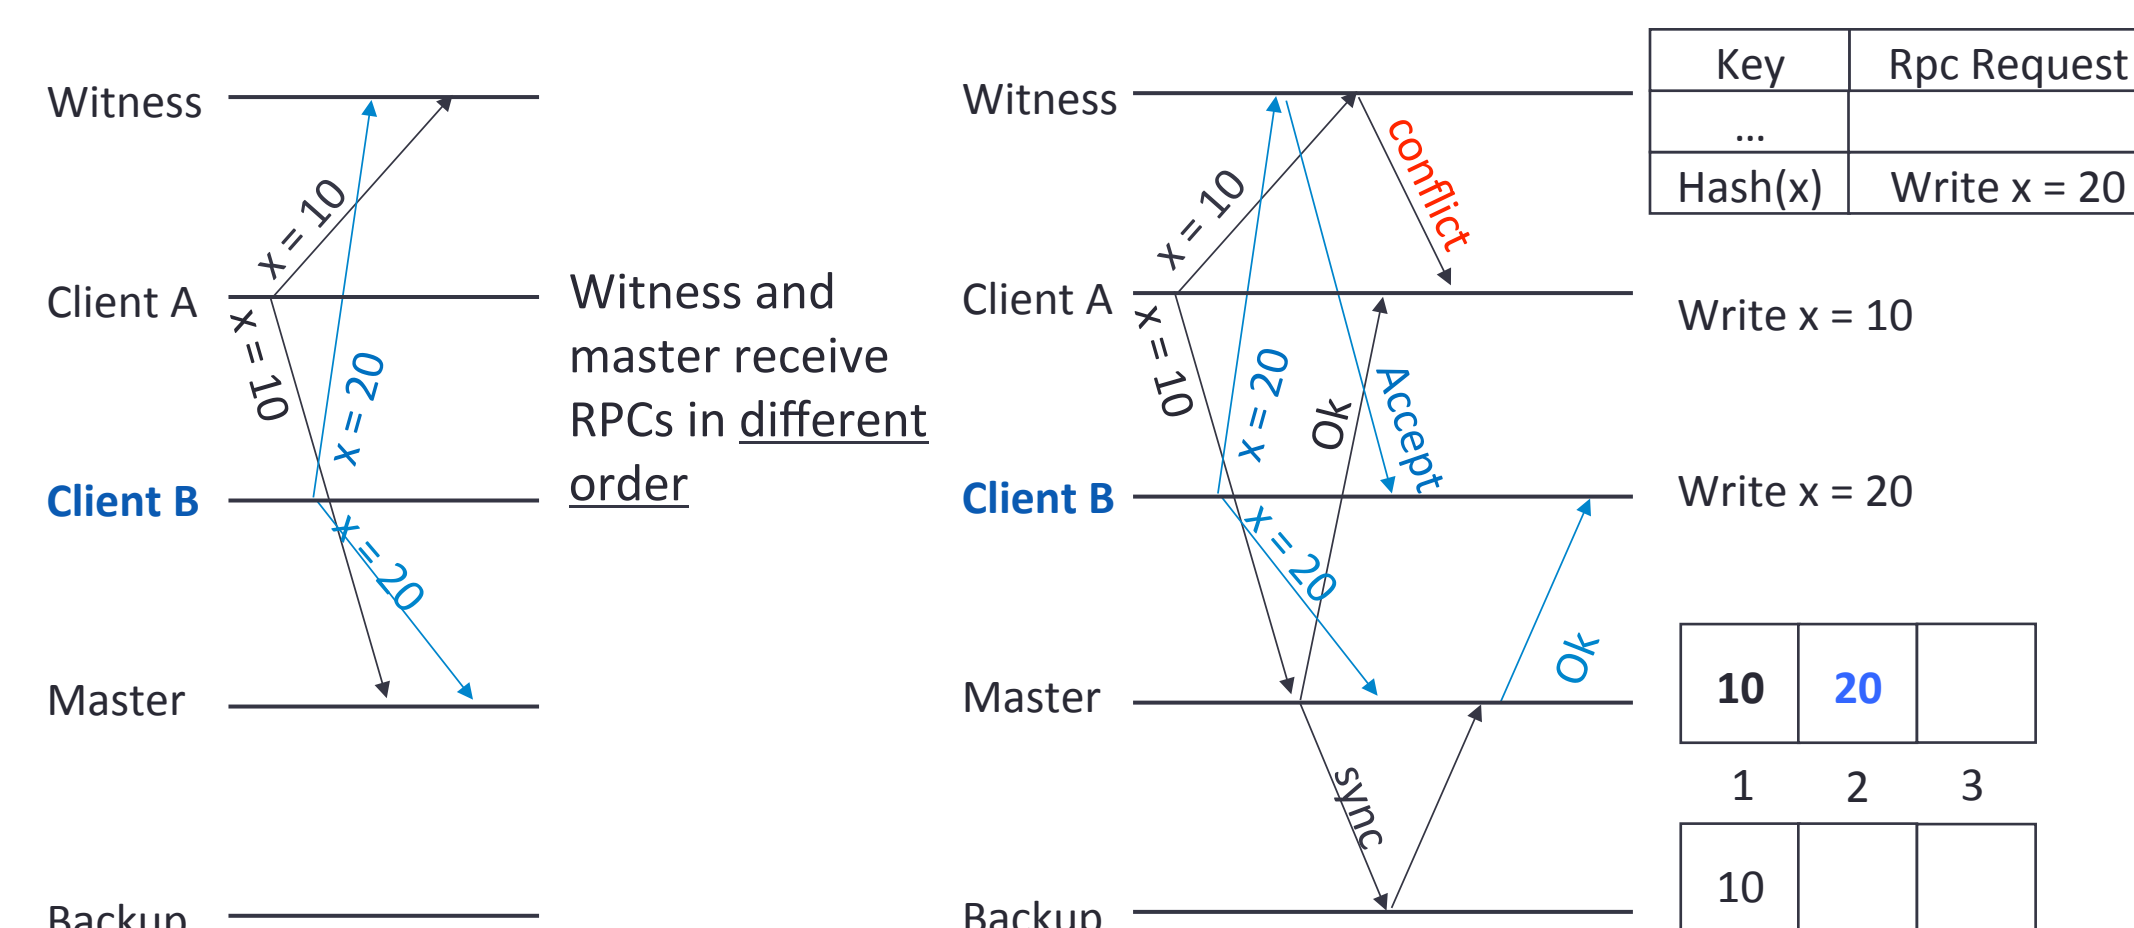
Applications

- Things don't care durability
- Split of update / validate client
- Many updates before externalization.

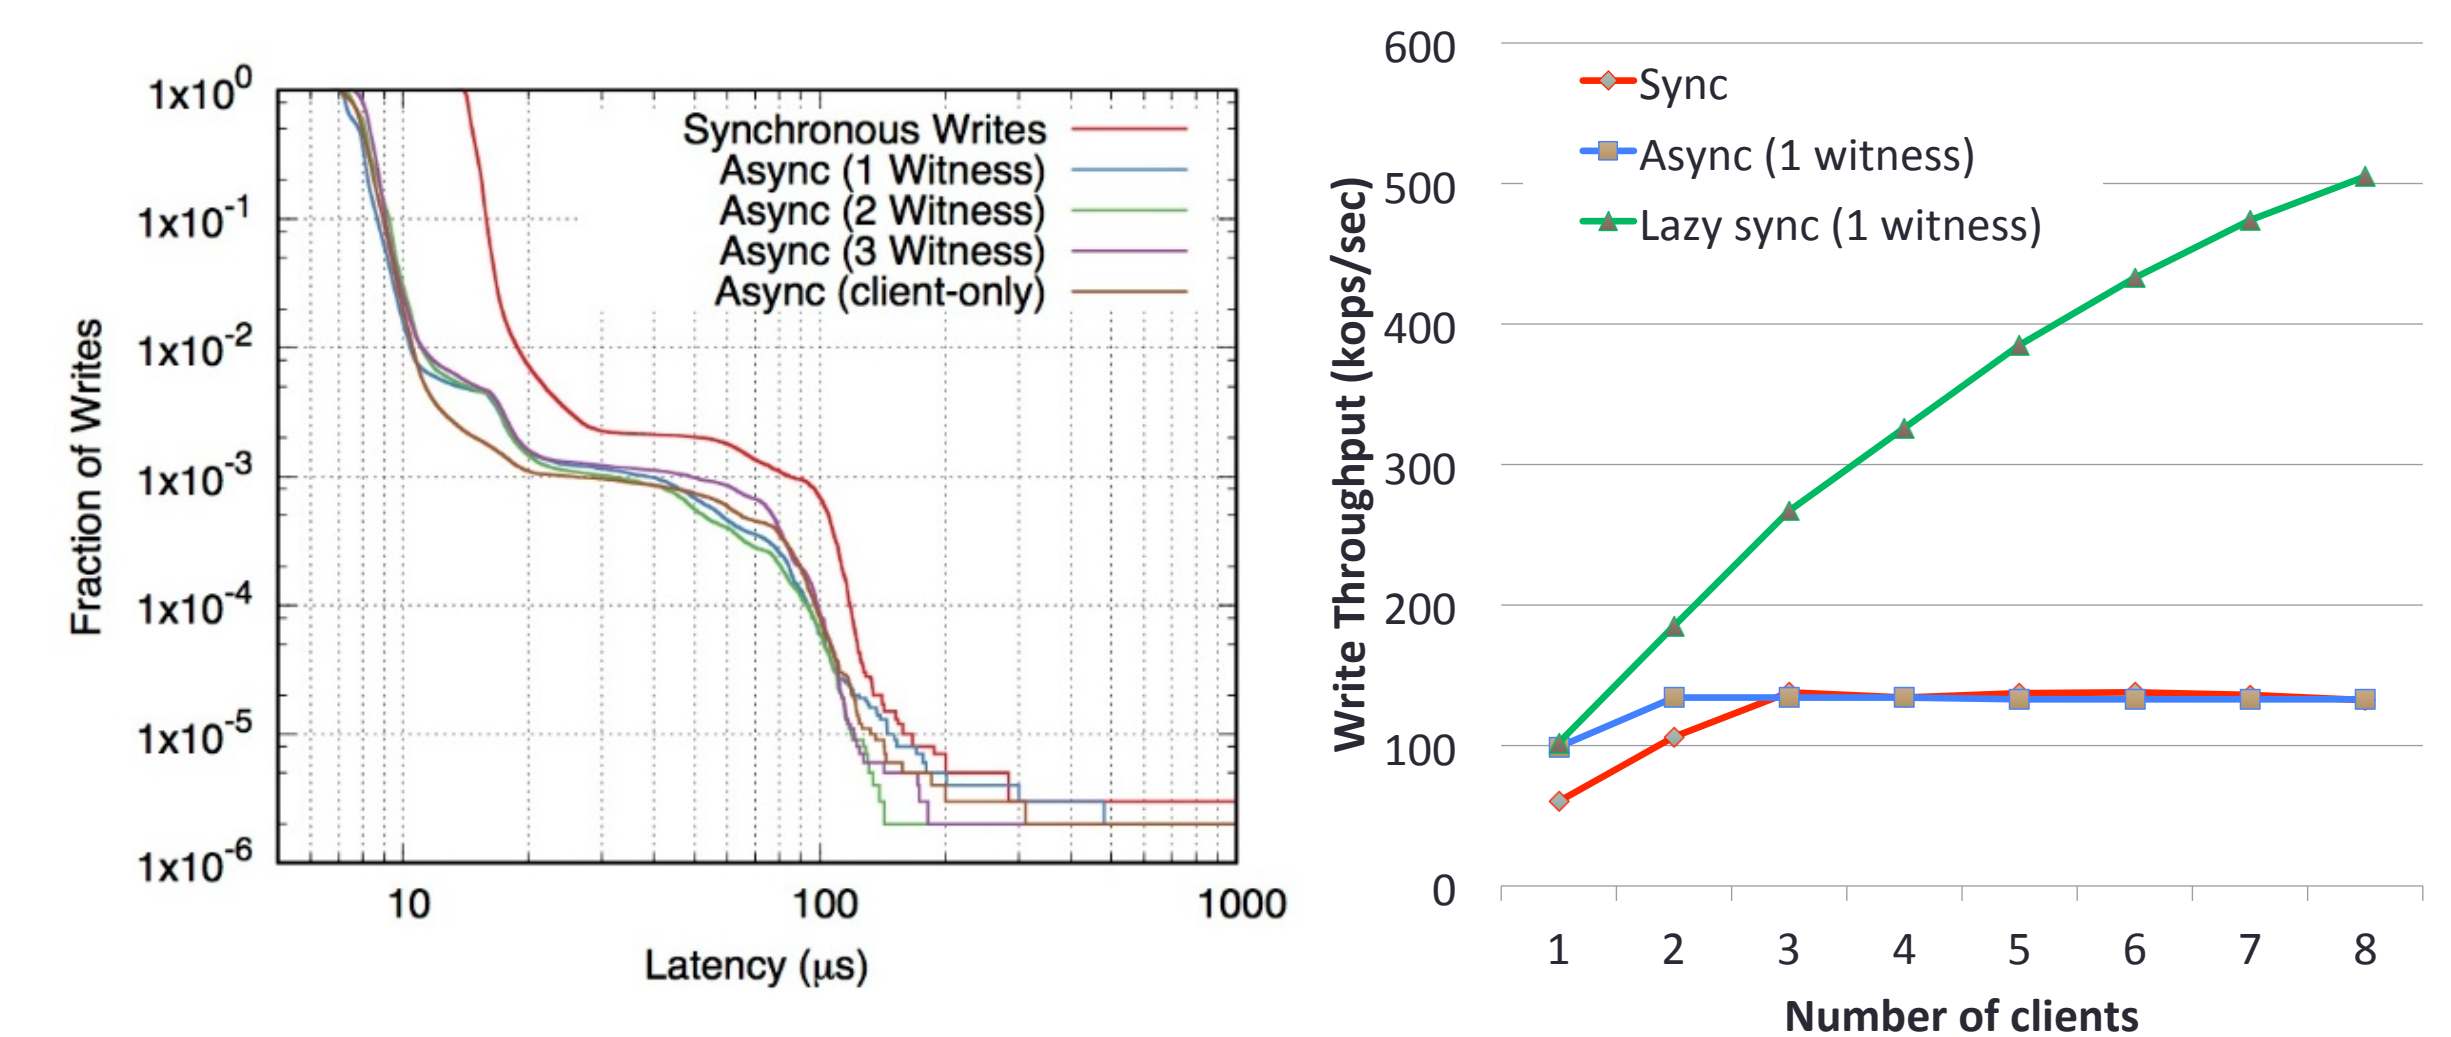

3RD PARTY-DRIVEN RECOVERY



Witness detects out-of-order recording and reject.



FASTER + HIGHER THROUGHPUT



CONCLUSIONS

- Retry RPC requests if a server crashes
 - \rightarrow strong consistency & durability (for witness) without synchronous backup.
- Decoupling durability process from critical path improved performance
 - Lower latency: 7.5 μ s regardless of backup media
 - Higher throughput: 137kops \rightarrow 500 kops
- RIFL (Reusable Infrastructure for Linearizability) eases design and reasoning of consistency